

Aprendizagem 2023
Homework II – Group 28

Gonalo Barias (ist1103124) & Raquel Braunschweig (ist1102624)

Part I: Pen and Paper

Consider the following dataset ($y_3 - y_5$ are all categorical variables and the domain of y_2 is $[0, 1]$):

D	y_1	y_2	y_3	y_4	y_5	y_6
\mathbf{x}_1	0.24	0.36	1	1	0	A
\mathbf{x}_2	0.16	0.48	1	0	1	A
\mathbf{x}_3	0.32	0.72	0	1	2	A
\mathbf{x}_4	0.54	0.11	0	0	1	B
\mathbf{x}_5	0.66	0.39	0	0	0	B
\mathbf{x}_6	0.76	0.28	1	0	2	B
\mathbf{x}_7	0.41	0.53	0	1	1	B
\mathbf{x}_8	0.38	0.52	0	1	0	A
\mathbf{x}_9	0.42	0.59	0	1	1	B

1. Consider $x_1 - x_7$ to be training observations, $x_8 - x_9$ to be testing observations, $y_1 - y_5$ to be input variables and y_6 to be the target variable.

Hint: you can use `scipy.stats.multivariate_normal` for multivariate distribution calculus

- (a) Learn a Bayesian classifier assuming: i) $\{y_1, y_2\}$, $\{y_3, y_4\}$ and $\{y_5\}$ sets of independent variables (e.g., $y_1 \perp\!\!\!\perp y_3$ yet $y_1 \not\perp\!\!\!\perp y_2$), and ii) $y_1 \times y_2 \in \mathbb{R}^2$ is normally distributed. Show all parameters (distributions and priors for subsequent testing).

Gonalo

- (b) Under a MAP assumption, classify each testing observation showing all your calculus.

Gonalo

- (c) Consider that the default decision threshold of $\theta = 0.5$ can be adjusted according to

$$f(\mathbf{x}|\theta) = \begin{cases} A, & P(A|\mathbf{x}) > \theta \\ B, & \text{otherwise} \end{cases}$$

Under a maximum likelihood assumption, what thresholds optimize testing accuracy?

Raquel

2. Let y_1 be the target numeric variable, $y_2 - y_6$ be the input variables where y_2 is binarized under an equal-width (equal-range) discretization. For the evaluation of regressors, consider a 3-fold cross-validation over the full dataset ($x_1 - x_9$) without shuffling the observations.

- (a) Identify the observations and features per data fold after the binarization procedure.

Raquel

- (b) Consider a distance-weighted kNN with $k = 3$, Hamming distance (d), and $1/d$ weighting. Compute the MAE of this kNN regressor for the 1st iteration of the cross-validation (i.e. train observations have the lower indices).

Raquel

Part II: Programming and critical analysis

Considering the `column_diagnosis.arff` dataset available at the course webpage's homework tab. Using `sklearn`, apply a 10-fold stratified cross-validation with shuffling (`random_state=0`) for the assessment of predictive models along this section.

1. Compare the performance of kNN with $k = 5$ and Naïve Bayes with Gaussian assumption (consider all remaining parameters for each classifier as `sklearn`'s default):

- (a) Plot two boxplots with the fold accuracies for each classifier.

```
1 import matplotlib.pyplot as plt, pandas as pd
2 from scipy.io.arff import loadarff
3 from sklearn.model_selection import cross_val_score, StratifiedKFold
4 from sklearn.neighbors import KNeighborsClassifier
5 from sklearn.naive_bayes import GaussianNB
6 from scipy.stats import ttest_rel
7
8 # Read the ARFF file and prepare data
9 data = loadarff("./data/column_diagnosis.arff")
10 df = pd.DataFrame(data[0])
11 df["class"] = df["class"].str.decode("utf-8")
12 X, y = df.drop("class", axis=1), df["class"]
13
14 # Initialize classifiers
15 knn_classifier = KNeighborsClassifier(n_neighbors=5)
16 naive_bayes_classifier = GaussianNB()
17
18 # Define cross-validation strategy
19 cv = StratifiedKFold(n_splits=10, shuffle=True, random_state=0)
20
21 # Evaluate classifiers
22 knn_accuracies = cross_val_score(knn_classifier, X, y, cv=cv, scoring='
    accuracy')
23 naive_bayes_accuracies = cross_val_score(naive_bayes_classifier, X, y, cv=cv,
    scoring='accuracy')
24
25 # Plot boxplots
26 plt.boxplot([knn_accuracies, naive_bayes_accuracies], labels=['kNN', 'Naive
    Bayes'])
27 plt.title('Classifier Comparison')
28 plt.ylabel('Accuracy')
29 plt.show()
```

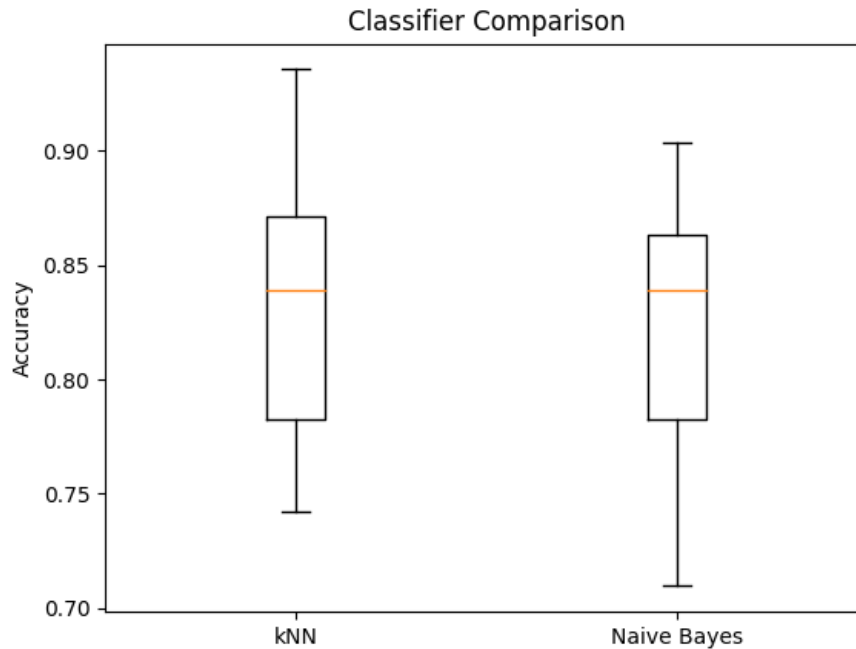


Figure 1: 10-Fold Cross-Validation Results for kNN and Naive Bayes on Column Diagnosis Dataset

- (b) Using **scipy**, test the hypothesis "kNN is statistically superior to Naïve Bayes regarding accuracy", asserting whether is true.

```

1 # Perform paired t-test
2 t_statistic, p_value = ttest_rel(knn_accuracies, naive_bayes_accuracies)
3
4 # Check the p-value
5 if p_value < 0.05:
6     print("Reject null hypothesis: kNN is statistically superior to Naive
7         Bayes in terms of accuracy")
8 else:
9     print("Fail to reject null hypothesis: No significant difference in
10         accuracy between kNN and Naive Bayes")

```

Answer: Fail to reject null hypothesis: No significant difference in accuracy between kNN and Naive Bayes

2. Consider two kNN predictors with $k = 1$ and $k = 5$ (uniform weights, Euclidean distance, all remaining parameters as default). Plot the differences between the two cumulative confusion matrices of the predictors. Comment.

```

1 import numpy as np, matplotlib.pyplot as plt, pandas as pd, seaborn as sns
2 from sklearn.model_selection import StratifiedKFold
3 from sklearn.neighbors import KNeighborsClassifier
4 from sklearn.metrics import confusion_matrix
5 from scipy.io.arff import loadarff
6
7 # Read the ARFF file and prepare data
8 data = loadarff("./data/column_diagnosis.arff")

```

```

9 df = pd.DataFrame(data[0])
10 df["class"] = df["class"].str.decode("utf-8")
11 X, y = df.drop("class", axis=1), df["class"]
12
13 # Initialize StratifiedKFold with 10 folds and shuffling
14 folds = StratifiedKFold(n_splits=10, shuffle=True, random_state=0)
15
16 # Create kNN classifiers with k=1 and k=5
17 knn_1 = KNeighborsClassifier(
18     n_neighbors=1, weights="uniform", metric="euclidean"
19 )
20 knn_5 = KNeighborsClassifier(
21     n_neighbors=5, weights="uniform", metric="euclidean"
22 )
23
24 labels = ["Hernia", "Normal", "Spondylolisthesis"]
25 cm_1, cm_5 = np.zeros((3, 3)), np.zeros((3, 3))
26 for train_k, test_k in folds.split(X, y):
27     X_train, X_test = X.iloc[train_k], X.iloc[test_k]
28     y_train, y_test = y.iloc[train_k], y.iloc[test_k]
29
30     # Fit kNN classifiers and assess
31     knn_1.fit(X_train, y_train)
32     knn_5.fit(X_train, y_train)
33     knn_1_pred, knn_5_pred = knn_1.predict(X_test), knn_5.predict(X_test)
34     cm_1 += np.array(confusion_matrix(y_test, knn_1_pred, labels=labels))
35     cm_5 += np.array(confusion_matrix(y_test, knn_5_pred, labels=labels))
36
37 # Calculate cumulative confusion matrices
38 cm_diff = cm_1 - cm_5
39 cm_diff_df = pd.DataFrame(cm_diff, index=labels, columns=labels)
40
41 # Plot the differences
42 plt.figure(figsize=(9, 7))
43 sns.heatmap(
44     cm_diff_df, cmap="Purples", annot=True, annot_kws={"fontsize": 14}, fmt="g"
45 )
46 plt.xlabel("Predicted")
47 plt.ylabel("Actual")
48 plt.show()

```




Figure 2: Confusion Matrix Differences Between k=1 and k=5 k-Nearest Neighbors (kNN) Classifiers

Blah

3. **Considering the unique properties of `column_diagnosis`, identify three possible difficulties of naïve Bayes when learning from the given dataset.**

Gonçalo