Aprendizagem 2023
Homework IV – Group 28

Gonçalo Bárias (ist1103124) & Raquel Braunschweig (ist1102624)

Part I: Pen and Paper

Given the following observations, $\left\{\begin{pmatrix}1\\0.6\\0.1\end{pmatrix},\begin{pmatrix}0\\-0.4\\0.8\end{pmatrix},\begin{pmatrix}0\\0.2\\0.5\end{pmatrix},\begin{pmatrix}1\\0.4\\-0.1\end{pmatrix}\right\}$.

Consider a Bayesian clustering that assumes $\{y_1\} \perp\!\!\!\perp \{y_2, y_3\}$, two clusters following a Bernoulli distribution on y_1 (p_1 and p_2), a multivariate Gaussian on $\{y_2, y_3\}$ (N_1 and N_2), and the following initial mixture:

$$\pi_1 = 0.5 \quad , \quad \pi_2 = 0.5$$

$$p_1 = P(y_1 = 1) = 0.3 \quad , \quad p_2 = P(y_1 = 1) = 0.7$$

$$\mathcal{N}_1\left(\boldsymbol{\mu}_1 = \begin{pmatrix}1\\1\end{pmatrix}, \boldsymbol{\Sigma}_1 = \begin{pmatrix}2 & 0.5\\0.5 & 2\end{pmatrix}\right) \quad , \quad \mathcal{N}_2\left(\boldsymbol{\mu}_2 = \begin{pmatrix}0\\0\end{pmatrix}, \boldsymbol{\Sigma}_2 = \begin{pmatrix}1.5 & 1\\1 & 1.5\end{pmatrix}\right)$$

1. **Perform one epoch of the EM clustering algorithm and determine the new parameters.**
   Hint: **we suggest you to use numpy and scipy, however disclose the intermediary results step by step.**

   The EM (Expectation-Maximization) algorithm has four major steps: Initialization, Expectation, Maximization and Evaluate.

   Initialization

   We'll start by labeling each observation:

$$x_1 = \begin{pmatrix}1\\0.6\\0.1\end{pmatrix} \quad , \quad x_2 = \begin{pmatrix}0\\-0.4\\0.8\end{pmatrix} \quad , \quad x_3 = \begin{pmatrix}0\\0.2\\0.5\end{pmatrix} \quad , \quad x_4 = \begin{pmatrix}1\\0.4\\-0.1\end{pmatrix}$$

   From the statement we have the following initial parameters, p_1, p_2, μ_1, μ_2, Σ_1, Σ_2, π_1 and π_2:

Cluster	p	μ	Σ	π
Cluster 1	0.3	$\begin{pmatrix}1\\1\end{pmatrix}$	$\begin{pmatrix}2 & 0.5\\0.5 & 2\end{pmatrix}$	0.5
Cluster 2	0.7	$\begin{pmatrix}0\\0\end{pmatrix}$	$\begin{pmatrix}1.5 & 1\\1 & 1.5\end{pmatrix}$	0.5

   Table 1: Initial parameters for the 2 clusters

   Expectation (E-step)

   Considering $\{y_1\} \perp\!\!\!\perp \{y_2, y_3\}$ we know the posterior probability, $P(c_k|x_i)$, is given by Baye's rule:

$$\gamma_{k,i} = P(c_k|x_i) = \frac{P(y_1, y_2, y_3|c_k)P(c_k)}{P(y_1, y_2, y_3)} = \frac{P(y_1|c_k)P(y_2, y_3|c_k)P(c_k)}{P(y_1)P(y_2, y_3)} \tag{1}$$

The variable y_1 follows a Bernoulli distribution ($y_1 \sim \text{Bern}\,(p = p_k)$), and so the likelihoods, $P(y_1 = 0|c_k)$ and $P(y_1 = 1|c_k)$, can be calculated for each cluster:

$$\begin{aligned} P(y_1 = 0|c_1) &= 1 - p_1 = 1 - 0.3 = 0.7 \qquad & P(y_1 = 0|c_2) &= 1 - p_2 = 1 - 0.7 = 0.3 \\ P(y_1 = 1|c_1) &= p_1 = 0.3 \qquad & P(y_1 = 1|c_2) &= p_2 = 0.7 \end{aligned}$$

We know the likelihood, $P(y_2, y_3|c_k)$, follows a multivariate Gaussian, and so it is given by (considering $d = 2$, since we are working in two dimensions):

$$P(y_2 = a, y_3 = b|c_k) = \mathcal{N}_k(y_2, y_3|\mu_k, \Sigma_k) = \frac{\exp\left(-\frac{1}{2}\left(\begin{bmatrix} a \\ b \end{bmatrix} - \mu_k\right)^T \Sigma_k^{-1} \left(\begin{bmatrix} a \\ b \end{bmatrix} - \mu_k\right)\right)}{(2\pi)^{d/2} \times |\Sigma_k|^{1/2}} \tag{2}$$

We now have all the building blocks to calculate the posterior probabilities for each combination of observation, x_i and cluster, c_k.

Maximization (M-step)

For each cluster, c_k, we will calculate the following in order to update the parameters:

$$N_k = \sum_i \gamma_{k,i}$$

$$p'_k = \frac{1}{N_k} \sum_i \gamma_{k,i} \cdot x_{i\,[y_1]}$$

$$\mu'_k = \frac{1}{N_k} \sum_i \gamma_{k,i} \cdot x_{i\,[y_2 \wedge y_3]}$$

$$\Sigma'_k = \frac{1}{N_k} \sum_i \gamma_{k,i} \cdot \left(x_{i\,[y_2 \wedge y_3]} - \mu'_k\right) \cdot \left(x_{i\,[y_2 \wedge y_3]} - \mu'_k\right)^T$$

Considering $N = \sum_k N_k$, we can also update the priors:

$$\pi'_k = \frac{N_k}{N}$$

We can now update the values for both clusters using the previous equations:

$$\boxed{N_1} = \sum_i \gamma_{1,i} = \gamma_{1,1} + \gamma_{1,2} + \gamma_{1,3} + \gamma_{1,4} = 1.54467$$

$$\boxed{N_2} = \sum_i \gamma_{2,i} = \gamma_{2,1} + \gamma_{2,2} + \gamma_{2,3} + \gamma_{2,4} = 2.45533$$

Cluster 1

$$p'_1 = \frac{1}{N_1}\sum_i \gamma_{1,i}\cdot x_{i[y_1]} = \frac{\gamma_{1,1}\cdot 1 + \gamma_{1,2}\cdot 0 + \gamma_{1,3}\cdot 0 + \gamma_{1,4}\cdot 1}{1.54467} = 0.23404$$

$$\mu'_1 = \frac{1}{N_1}\sum_i \gamma_{1,i}\cdot x_{i[y_2\wedge y_3]} = \frac{\gamma_{1,1}\cdot\begin{pmatrix}0.6\\0.1\end{pmatrix} + \gamma_{1,2}\cdot\begin{pmatrix}-0.4\\0.8\end{pmatrix} + \gamma_{1,3}\cdot\begin{pmatrix}0.2\\0.5\end{pmatrix} + \gamma_{1,4}\cdot\begin{pmatrix}0.4\\-0.1\end{pmatrix}}{1.54467} = \begin{pmatrix}0.02651\\0.50713\end{pmatrix}$$

$$\begin{aligned}
\Sigma'_1 &= \frac{1}{N_1}\sum_i \gamma_{1,i}\cdot\left(x_{i[y_2\wedge y_3]} - \mu'_1\right)\cdot\left(x_{i[y_2\wedge y_3]} - \mu'_1\right)^T \\
&= \frac{1}{1.54467}\times\Bigg[\gamma_{1,1}\cdot\left(\begin{pmatrix}0.6\\0.1\end{pmatrix} - \begin{pmatrix}0.02651\\0.50713\end{pmatrix}\right)\cdot\left(\begin{pmatrix}0.6\\0.1\end{pmatrix} - \begin{pmatrix}0.02651\\0.50713\end{pmatrix}\right)^T \\
&+ \gamma_{1,2}\cdot\left(\begin{pmatrix}-0.4\\0.8\end{pmatrix} - \begin{pmatrix}0.02651\\0.50713\end{pmatrix}\right)\cdot\left(\begin{pmatrix}-0.4\\0.8\end{pmatrix} - \begin{pmatrix}0.02651\\0.50713\end{pmatrix}\right)^T \\
&+ \gamma_{1,3}\cdot\left(\begin{pmatrix}0.2\\0.5\end{pmatrix} - \begin{pmatrix}0.02651\\0.50713\end{pmatrix}\right)\cdot\left(\begin{pmatrix}0.2\\0.5\end{pmatrix} - \begin{pmatrix}0.02651\\0.50713\end{pmatrix}\right)^T \\
&+ \gamma_{1,4}\cdot\left(\begin{pmatrix}0.4\\-0.1\end{pmatrix} - \begin{pmatrix}0.02651\\0.50713\end{pmatrix}\right)\cdot\left(\begin{pmatrix}0.4\\-0.1\end{pmatrix} - \begin{pmatrix}0.02651\\0.50713\end{pmatrix}\right)^T\Bigg] \\
&= \begin{pmatrix}0.14137 & -0.10541\\-0.10541 & 0.09605\end{pmatrix}
\end{aligned}$$

$$\pi'_1 = \frac{N_1}{N} = \frac{N_1}{N_1+N_2} = \frac{1.54467}{1.54467+2.45533} = 0.38617$$

Cluster 2

$$p'_2 = \frac{1}{N_2}\sum_i \gamma_{2,i}\cdot x_{i[y_1]} = \frac{\gamma_{2,1}\cdot 1 + \gamma_{2,2}\cdot 0 + \gamma_{2,3}\cdot 0 + \gamma_{2,4}\cdot 1}{2.45533} = 0.66732$$

$$\mu'_2 = \frac{1}{N_2}\sum_i \gamma_{2,i}\cdot x_{i[y_2\wedge y_3]} = \frac{\gamma_{2,1}\cdot\begin{pmatrix}0.6\\0.1\end{pmatrix} + \gamma_{2,2}\cdot\begin{pmatrix}-0.4\\0.8\end{pmatrix} + \gamma_{2,3}\cdot\begin{pmatrix}0.2\\0.5\end{pmatrix} + \gamma_{2,4}\cdot\begin{pmatrix}0.4\\-0.1\end{pmatrix}}{2.45533} = \begin{pmatrix}0.30914\\0.21042\end{pmatrix}$$

$$\Sigma_2' = \frac{1}{N_2} \sum_i \gamma_{2,i} \cdot \left(x_{i\,[y_2 \wedge y_3]} - \mu_2'\right) \cdot \left(x_{i\,[y_2 \wedge y_3]} - \mu_2'\right)^T$$

$$= \frac{1}{2.45533} \times \Bigg[\gamma_{2,1} \cdot \left(\begin{pmatrix} 0.6 \\ 0.1 \end{pmatrix} - \begin{pmatrix} 0.30914 \\ 0.21042 \end{pmatrix} \right) \cdot \left(\begin{pmatrix} 0.6 \\ 0.1 \end{pmatrix} - \begin{pmatrix} 0.30914 \\ 0.21042 \end{pmatrix} \right)^T$$

$$+ \gamma_{2,2} \cdot \left(\begin{pmatrix} -0.4 \\ 0.8 \end{pmatrix} - \begin{pmatrix} 0.30914 \\ 0.21042 \end{pmatrix} \right) \cdot \left(\begin{pmatrix} -0.4 \\ 0.8 \end{pmatrix} - \begin{pmatrix} 0.30914 \\ 0.21042 \end{pmatrix} \right)^T$$

$$+ \gamma_{2,3} \cdot \left(\begin{pmatrix} 0.2 \\ 0.5 \end{pmatrix} - \begin{pmatrix} 0.30914 \\ 0.21042 \end{pmatrix} \right) \cdot \left(\begin{pmatrix} 0.2 \\ 0.5 \end{pmatrix} - \begin{pmatrix} 0.30914 \\ 0.21042 \end{pmatrix} \right)^T$$

$$+ \gamma_{2,4} \cdot \left(\begin{pmatrix} 0.4 \\ -0.1 \end{pmatrix} - \begin{pmatrix} 0.30914 \\ 0.21042 \end{pmatrix} \right) \cdot \left(\begin{pmatrix} 0.4 \\ -0.1 \end{pmatrix} - \begin{pmatrix} 0.30914 \\ 0.21042 \end{pmatrix} \right)^T \Bigg]$$

$$= \begin{pmatrix} 0.10829 & -0.08865 \\ -0.08865 & 0.10412 \end{pmatrix}$$

$$\pi_2' = \frac{N_2}{N} = \frac{N_2}{N_1 + N_2} = \frac{2.45533}{1.54467 + 2.45533} = 0.61383$$

Evaluate the log likelihood

Since we are only performing one epoch of the EM clustering algorithm, we can skip this step.

Conclusion

After performing one epoch of the EM clustering algorithm, we end up with the following updated parameters for each cluster:

Cluster	p'	μ'	Σ'	π'
Cluster 1	0.23404	$\begin{pmatrix} 0.02651 \\ 0.50713 \end{pmatrix}$	$\begin{pmatrix} 0.14137 & -0.10541 \\ -0.10541 & 0.09605 \end{pmatrix}$	0.38617
Cluster 2	0.66732	$\begin{pmatrix} 0.30914 \\ 0.21042 \end{pmatrix}$	$\begin{pmatrix} 0.10829 & -0.08865 \\ -0.08865 & 0.10412 \end{pmatrix}$	0.61383

Table 2: Updated parameters for the 2 clusters

2. **Given the new observation,** $x_{new} = \begin{bmatrix} 1 & 0.3 & 0.7 \end{bmatrix}^T$**, determine the cluster memberships (posteriors).**

   Raquel

3. **Performing a hard assignment of observations to clusters under a ML assumption, identify the silhouette of both clusters under a Manhattan distance.**

   Raquel

4. **Knowing the purity of the clustering solution is 0.75, identify the number of possible classes (ground truth).**

   Raquel

Part II: Programming and critical analysis

Recall the `column_diagnosis.arff` dataset from previous homeworks. For the following exercises, normalize the data using sklearn's `MinMaxScaler`.

1. **Using `sklearn`, apply *k*-means clustering fully unsupervisedly on the normalized data with** $k \in \{2, 3, 4, 5\}$ (random = 0 **and remaining parameters as default). Assess the silhouette and purity of the produced solutions.**

   Gonçalo

2. **Consider the application of PCA after the data normalization:**

   (a) **Identify the variability explained by the top two principal components.**

   Raquel

   (b) **For each one of these two components, sort the input variables by relevance by inspecting the absolute weights of the linear projection.**

   Raquel

3. **Visualize side-by-side the data using: i) the ground diagnoses, and ii) the *previously* learned** $k = 3$ **clustering solution. To this end, projected the normalized data onto a 2-dimensional data space using PCA and then color observations using the reference and cluster annotations.**

   Raquel

4. **Considering the results from questions (1) and (3), identify two ways on how clustering can be used to characterize the population of ill and healthy individuals.**

   Raquel